


FLOATING POINT NUMBERS



Assignment Project Exam Help

IEEE floating point standard

<https://tutorcs.com>

WeChat: cstutorcs

Bernhard Kainz (with thanks to **A. Gopalan**, **N. Dulay** and **E. Edwards**)

b.kainz@imperial.ac.uk

IEEE floating point standard


- IEEE: institute of electrical and electronic engineers (USA)
- Comprehensive standard for binary floating point arithmetic
- Widely adopted → predictable results independent of architecture
- Standard defines:
 - **Format** of binary floating point numbers, i.e. how the fields are stored in memory
 - **Semantics** of arithmetic operations
 - Rules for **error conditions**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Single precision format (32-bit)



- Coefficient is called the **significand** in the IEEE standard
- Value represented is $\pm 1.F \times 2^{E-127}$
- The **normal bit** (the 1.) is omitted from the significand field → a **hidden bit**
- Single precision yields **24 bits** (approx. **7 decimal digits** of precision)
- **Normalised ranges** in decimal are approximately:
 -10^{38} to -10^{-38} , **0**, 10^{38} to 10^{-38}

Exponent field

- In the IEEE standard, exponents are stored as **excess values**, not as 2's complement

Assignment Project Exam Help

- Example: **In 8-bit excess-127**

<https://tutorcs.com>
-127 would be held as 0000 0000

WeChat: cstutorcs
0 0111 1111

1 1000 0000

...
128 1111 1111

- Allows non-negative floating point numbers to be compared using simple integer comparisons

Double precision format (64-bit)



- Value represented is $\pm 1.F \times 2^E$
- Double precision yields **53 bits** (approx. **16 decimal digits** of precision)
- **Normalised ranges** in decimal are approximately:
 -10^{308} to -10^{-308} , **0**, 10^{308} to 10^{-308}
- Single precision generally reserved for when memory is scarce or for debugging numerical calculations since rounding errors show up more quickly

Example: conversion to IEEE format

What is 42.6875 in IEEE single precision format?


1. Convert to binary number: $42.6875 = 101010.1011$
2. **Normalise:** $1.0101\ 0101\ 1 \times 2^5$
3. **Significand field** is thus:
 $01010101\ 1000\ 0000\ 0000\ 000$
4. **Exponent field** is $(5 + 127 = 132)$: $1000\ 0100$

Sign S	Exponent E	Significand F
0	1000 0100	0101 0101 1000 0000 0000 000

Hex: 422A C000

Example: conversion from IEEE format

What is the IEEE single precision value represented by **BEC0 0000** in decimal?



1. **Exponent field:** 0111 1101 = 125
2. **True binary exponent:** $125 - 127 = -2$
3. **Significand field + hidden bit:**
 $1.1000\ 0000\ 0000\ 0000\ 0000\ 000$
4. **So unsigned value** is $1.1 \times 2^{-2} = 0.011$ (binary)
 $= 0.25 + 0.125 = 0.375$ (decimal)
5. Adding **sign bit** gives finally **-0.375**

Example: addition

Carry out the addition $42.6875 + 0.375$ in IEEE single precision arithmetic

Assignment Project Exam Help

Number	Sign	Exponent	Significand
42.6875	0	1000 0100	0101 0101 1000 0000 0000 000
0.375	0	0111 1101	1000 0000 0000 0000 0000 000

WeChat: cstutorcs

- To add these numbers, exponents must be the same → make the smaller exponent equal to the larger by shifting significand accordingly
- **Note:** must restore **hidden bit** when carrying out floating point operations

Example: addition (cont.)

- **Significand** of larger no.: 1.0101 0101 1000 0000 0000 000
- **Significand** of smaller no.: 1.1000 0000 0000 0000 0000 000

Assignment Project Exam Help

- Exponents differ by $(1000\ 0100 - 0111\ 1101 = 7)$ so shift binary point of smaller no. 7 places to the left:

<https://tutorcs.com>

- **Significand** of smaller no.: 0.0000 0011 0000 0000 0000 000
- **Significand** of larger no.: 1.0101 0101 1000 0000 0000 000
- **Significand** of **sum**: 1.0101 1000 1000 0000 0000 000

WeChat: cstutorcs

- So **sum** is $1.0101\ 1000\ 1 \times 2^5 = 10\ 1011.0001 = 43.0625$

Sign S	Exponent E	Significand F
0	1000 0100	0101 1000 1000 0000 0000 000

Special values

- IEEE formats can encode five kinds of values: **zero**, **normalised numbers**, **denormalised numbers**, **infinity** and **not-a-number (NaNs)**
- Single precision representations:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

IEEE value	Sign field	Exponent	Significand	True exponent
± 0	0 or 1	0	0 (all zeros)	
\pm denormalised no.	0 or 1	0	Any non-zero bit pattern	-126
\pm normalised no.	0 or 1	1 ... 254	Any bit pattern	-126 ... 127
$\pm \infty$	0 or 1	255	0 (all zeros)	
Not-a-number	0 or 1	255	Any non-zero bit pattern	

Denormalised numbers

- An **all zero exponent** is used to represent both **zero** and **denormalised numbers**
- An **all one exponent** is used to represent **infinities** and **not-a-numbers**
- Means **range for normalised numbers is reduced**, for single precision the exponent range is $-126 \dots 127$ rather than $-127 \dots 128$
- **Denormalised numbers** represent values between the underflow limits and zero, i.e. for single precision we have $\pm 0.F \times 2^{-126}$
- Allows a more **gradual shift to zero** – useful in some numerical applications

Infinites and NaNs

- Infinites represent values **exceeding the overflow limits** and for divisions of non-zero quantities by zero
- You can do basic 'arithmetic' with them, e.g.:

$$\infty + 5 = \infty, \quad \infty + \infty = \infty$$

- NaNs represent the result of operations which have **no (real) mathematical interpretation**, e.g.

$$\frac{0}{0}, \quad +\infty + -\infty, \quad 0 \times \infty, \quad \text{square root of a negative number}$$

- Operations resulting in NaNs can either yield a NaN result (**quiet NaN**) or an exception (**signalling NaN**)

Special Operations

Operation	Result
$N \div \pm \text{Infinity}$	0
$\pm \text{Infinity} \times \pm \text{Infinity}$	$\pm \text{Infinity}$
$\pm \text{non-zero} \div 0$	$\pm \text{Infinity}$
$\text{Infinity} + \text{Infinity}$	Infinity
$\pm 0 \div \pm 0$	<i>NaN</i>
$\text{Infinity} - \text{Infinity}$	<i>NaN</i>
$\pm \text{Infinity} \div \pm \text{Infinity}$	<i>NaN</i>
$\pm \text{Infinity} \times 0$	<i>NaN</i>

© Assignment Project Exam Help ©
SOME FUN

<https://tutorcs.com>

WeChat: cstutorcs

Floating Point Precision

- C code:

```
#include <stdio.h>
```

```
int main() {
```

```
    float a, b, c;
```

```
    float EPSILON = 0.0000001;
```

```
    a = 1.345f; b = 1.123f;
```

```
    c = a + b;
```

```
    if (c == 2.468)
```

```
        printf ("They are equal.\n");
```

```
    else
```

```
        printf ("\nThey are not equal! The value of c is %.10f or %f\n",c,c);
```

```
    // With some tolerance
```

```
    if (((2.468 - EPSILON) < c) && (c < (2.468 + EPSILON)))
```

```
        printf ("\n%.10f is equal to 2.468 with tolerance\n\n", c);
```

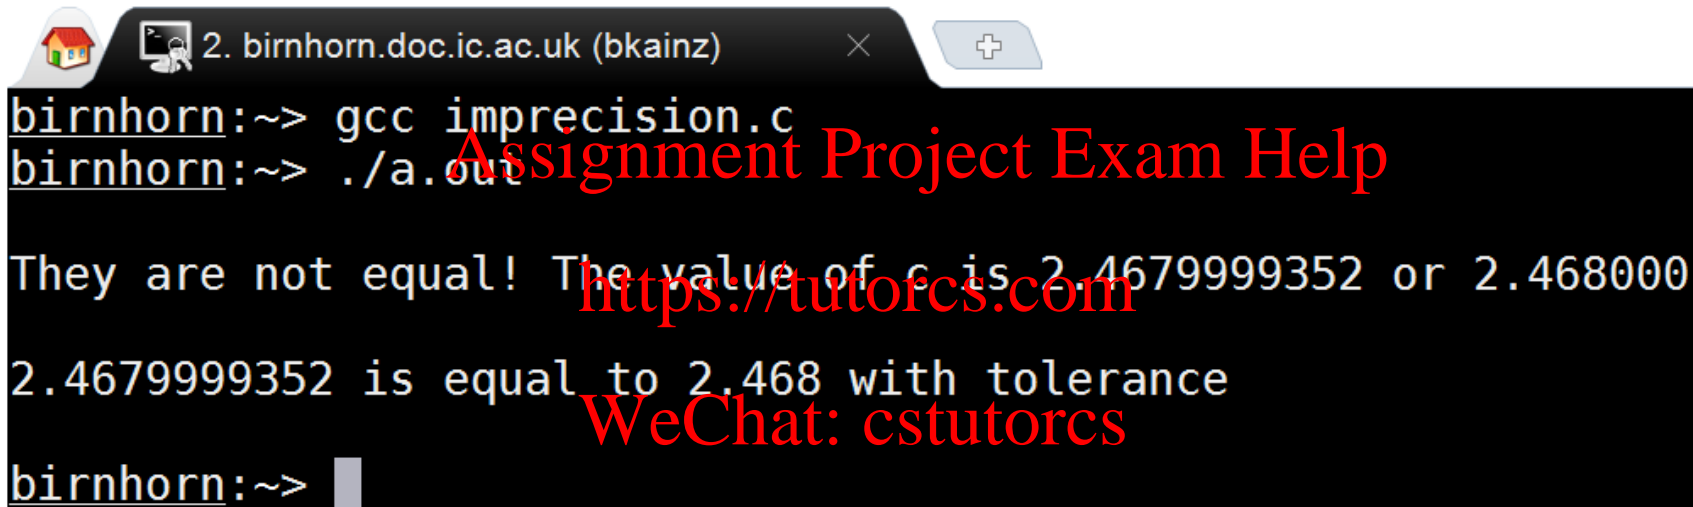
```
}
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Run-time



A terminal window with a dark background and white text. The window has a title bar with a home icon, a terminal icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". The terminal shows the following commands and output:

```
birnhorn:~> gcc imprecision.c
birnhorn:~> ./a.out

They are not equal! The value of c is 2.4679999352 or 2.468000

2.4679999352 is equal to 2.468 with tolerance

birnhorn:~> █
```

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

Finding Machine Epsilon

- Pseudo-code

Set machineEps = 1.0,

Loop <https://tutorcs.com>

machineEps = machineEps/2.0
WeChat: cstutorcs

Until $((1 + \text{machineEps}/2.0) \neq 1)$

Print machineEps

Finding Machine Epsilon

- C code

```
#include <stdio.h>
```

```
int main( int argc, char **argv )
```

```
{
```

```
    float machEps = 1.0f;
```

```
    do {
```

```
        machEps /= 2.0f;
```

```
        // If next epsilon yields 1, then break, because current
```

```
        // epsilon is the machine epsilon.
```

```
    }
```

```
    while ((float)(1.0 + (machEps/2.0f)) != 1.0);
```

```
    printf( "\nCalculated Machine epsilon: %G\n\n", machEps );
```

```
    return 0;
```

```
}
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Finding Machine Epsilon

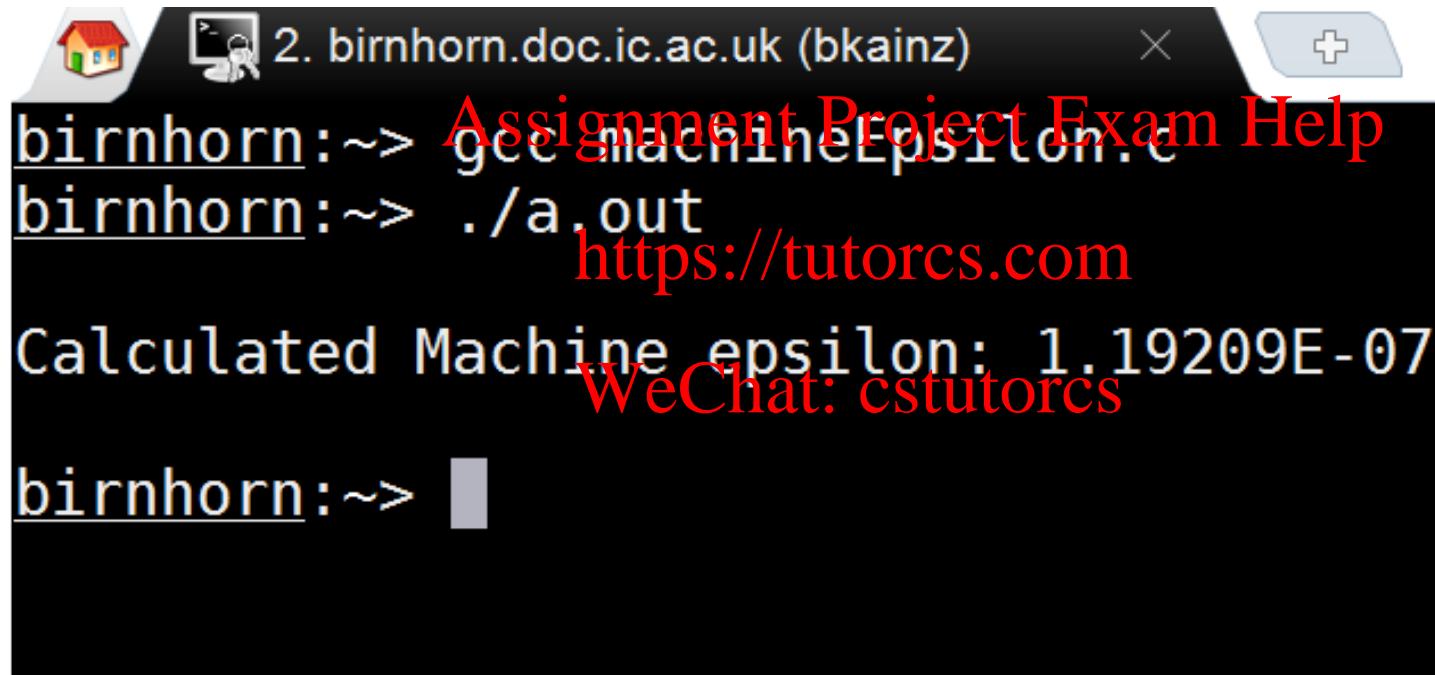
- In Java

```
public class machEps
{
    private static void calculateMachineEpsilonFloat() {
        float machEps = 1.0f;
        do {
            machEps /= 2.0f;
        } while ((float)(1.0 - (machEps / 2.0)) != 1.0);

        System.out.println( "Calculated machine epsilon: " + machEps );
    }

    public static void main (String args[])
    {
        calculateMachineEpsilonFloat ();
    }
}
```

Run-time



A terminal window with a dark background and light-colored text. The window title bar shows a home icon, a terminal icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". The terminal content shows a user at the "birnhorn" prompt running a C compiler and then a program. The program outputs the calculated machine epsilon. There are red watermarks overlaid on the terminal text.

```
birnhorn:~> gcc machineEpsilon.c  
birnhorn:~> ./a.out  
Calculated Machine epsilon: 1.19209E-07  
birnhorn:~> 
```

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

Special Operations

- Example

```
#include <stdio.h>
```

```
int main (int argc, char **argv)  
{
```

```
    float a = 100.0;
```

```
    float b = a * -100;
```

```
    float c = b/a;
```

```
    int d = 2 * 10 + 3;
```

```
    printf ("\nValue of a = %f\n\n", a);
```

```
    printf ("\nValue of b = %f\n\n", b);
```

```
    printf ("\nValue of c = %f\n\n", c);
```

```
    printf ("\nValue of d = %d\n\n", d);
```

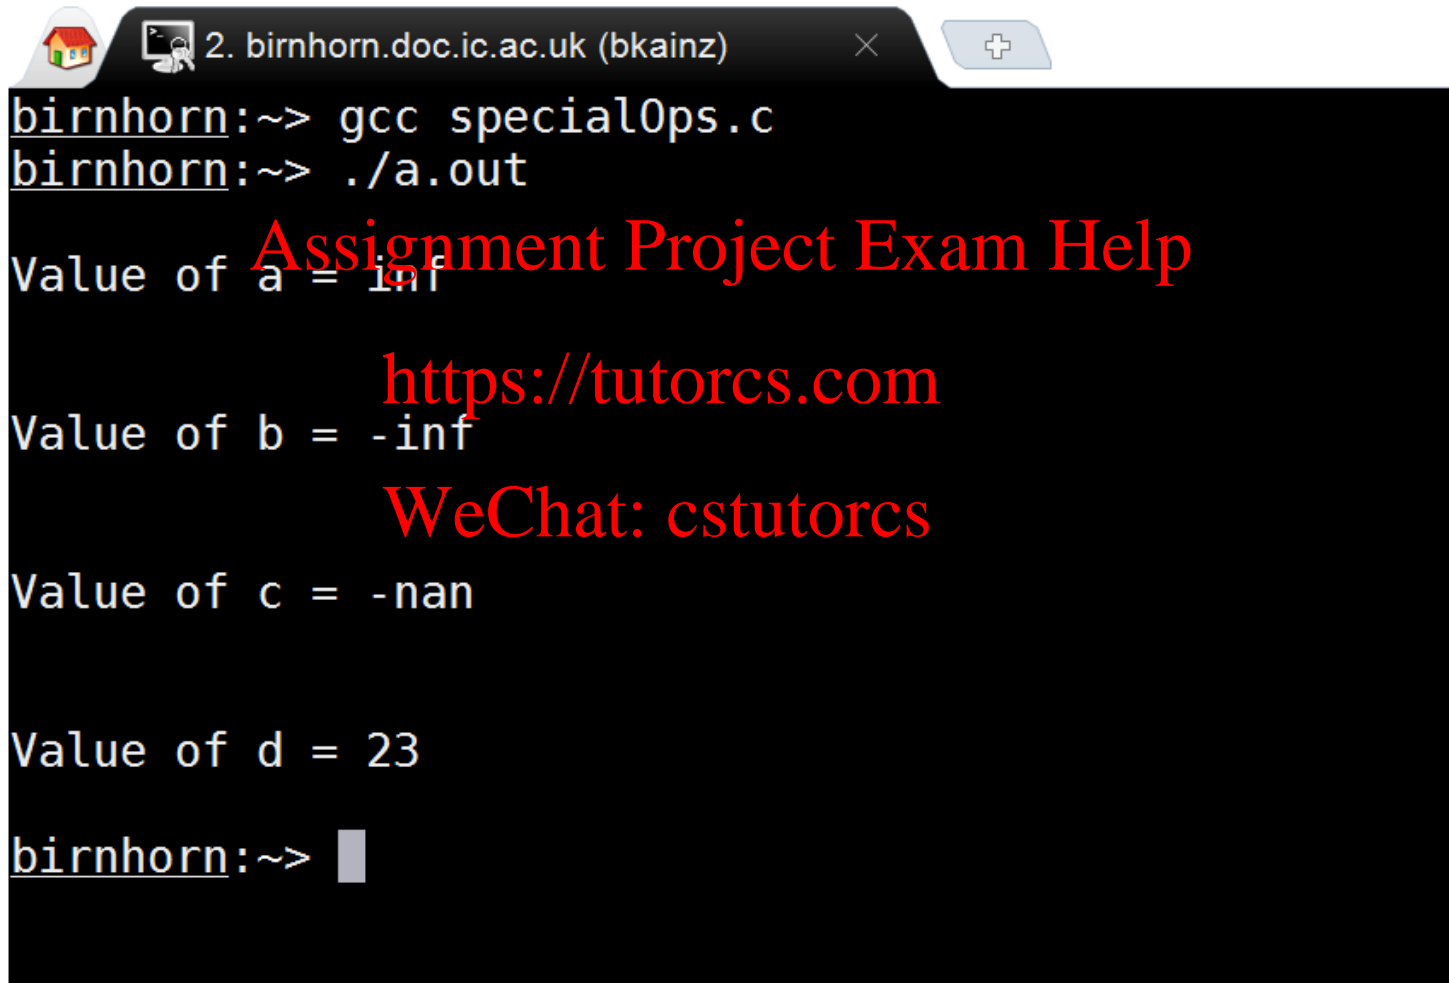
```
}
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Run-time

A terminal window with a dark background. The title bar shows a home icon, a document icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". The terminal content shows a user named "birnhorn" at a shell prompt. They run "gcc special0ps.c" and then "./a.out". The program outputs four lines: "Value of a = inf", "Value of b = -inf", "Value of c = -nan", and "Value of d = 23". The prompt "birnhorn:~>" is followed by a cursor.

```
2. birnhorn.doc.ic.ac.uk (bkainz) × +
birnhorn:~> gcc special0ps.c
birnhorn:~> ./a.out
Value of a = inf
Value of b = -inf
Value of c = -nan
Value of d = 23
birnhorn:~> █
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs